

Lyft_Baywheels_capstone_2019

August 30, 2019

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [319]: import certifi
import ssl
import geopy.geocoders
ctx = ssl.create_default_context(cafile=certifi.where())
geopy.geocoders.options.default_ssl_context = ctx
from geopy.geocoders import GoogleV3

In [463]: def get_file_contents(filename):
    """ Given a filename,
        return the contents of that file
        """
    try:
        with open(filename, 'r') as f:
            # It's assumed our file contains a single line,
            # with our API key
            return f.read().strip()
    except FileNotFoundError:
        print("%s' file not found" % filename)

secret_api_key = get_file_contents('secret.py')
```

1 Dear user, please enter a San Francisco address below:

By entering an address and running the cells below, you will get a prediction of the average weekly rides that would be taken from your proposed station. You will also get a normalized number of rides predicted, for comparison to other stations. Basically, above 1.0 is more than average and below 1.0 is less than average. From that you get a recommendation of whether your proposed address is a good location for a new bike station. You will also see a map of the current stations and their usage and your proposed new location.

```
In [464]: address = "1800 Montgomery St, San Francisco, CA 94111"
#example: address = "1800 Montgomery St, San Francisco, CA 94111"
```

```
In [502]: geolocator = GoogleV3(api_key=secret_api_key)
location = geolocator.geocode(address, timeout=10)
print(location.address)
print((location.latitude, location.longitude))
proposed_station = [location.latitude, location.longitude]
print('Calculating...')
coordinate_predictions = coordinates_to_predictions(proposed_station)
```

1800 Montgomery St, San Francisco, CA 94111, USA

(37.80599309999999, -122.4051806)

Calculating...

The predicted number of trips per week, starting at this location, is: 612

The normalized weekly trip count is: 3.58556842726421

The trip count is more than average, this may be a good location.

```
In [496]: import folium
```

```
plt.figure(figsize=(15,10))
m = folium.Map(location=[37.786375, -122.404904], tiles='Stamen Toner', zoom_start=13)

mapping_data.apply(lambda row:folium.CircleMarker(location=[row["start_station_latitude"],
row["start_station_longitude"]], radius=7,\
color=row['marker_color'], fill=True, fill_opacity=0.8\
).add_to(m), axis=1)
legend_html = '''
    <div style="position: fixed;
        bottom: 50px; right: 50px; width: 150px; height: 130px;
        border:2px solid grey; z-index:9999; font-size:14px;
        background-color:lightgrey;
        ">&nbsp;   <b>Number of Trips:</b><br>
        &nbsp;   <i class="fa fa-circle" style="color:yellow"></i>
        &nbsp;   <i class="fa fa-circle" style="color:orange"></i>
        &nbsp;   <i class="fa fa-circle" style="color:red"></i>
        &nbsp;   <i class="fa fa-circle" style="color:darkred"></i>
        &nbsp;   <i class="fa fa-circle" style="color:blue"></i>

    </div>
</div>
</div>
'''

m.get_root().html.add_child(folium.Element(legend_html))
m
```


```
Out[496]: <folium.folium.Map at 0x16e4a4470>
```

<Figure size 1080x720 with 0 Axes>

In case the html map does not render properly, here is a png snapshot of what it should look like:

```
In [477]: import matplotlib.image as mpimg

plt.figure(figsize = (100,20))
img=mpimg.imread('lyft_stations_screen_shot.png')
imgplot = plt.imshow(img)
```



1.0.1 That was the final product. Now let's back up and show how we got to this point.

2 1. Introduction and Motivation

Bike shares are great, but you may be wondering why there isn't a station near your house or a busy part of town. You may want to propose a new station. A simple way to evaluate whether your proposed station will be viable, and have enough users, is to compare to existing stations. Since bike share bikes are often used for getting to or from public transit, the distance to BART and Muni will also be considered.

The Lyft Baywheels bike share provides data for trips taken using their bikes. The data includes the station names, locations, and number of trips started and finished there. From this data I calculated the number of trips per station per week and found which stations are more or less popular. I also calculated how far the stations are from public transit. I fed this into a machine learning model to predict the popularity of a proposed station location.

3 2. Let's read in all the csv's and append them into one dataframe

3.0.1 The website for getting data is the Lyft Baywheels site:

<https://s3.amazonaws.com/baywheels-data/index.html>

```
In [5]: df_201907= pd.read_csv('201907-baywheels-tripdata.csv')
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/IPython/core/interactivity=interactivity, compiler=compiler, result=result)
```

```

In [6]: df_201906= pd.read_csv('201906-baywheels-tripdata.csv')
In [8]: df_201905= pd.read_csv('201905-baywheels-tripdata.csv')
In [9]: df_201904 = pd.read_csv('201904-fordgobike-tripdata.csv')
In [10]: df_201903 = pd.read_csv('201903-fordgobike-tripdata.csv')
In [11]: df_201902 = pd.read_csv('201902-fordgobike-tripdata.csv')
In [12]: df_201901 = pd.read_csv('201901-fordgobike-tripdata.csv')
In [13]: df_201812 = pd.read_csv('201812-fordgobike-tripdata.csv')
In [14]: df_201811 = pd.read_csv('201811-fordgobike-tripdata.csv')
In [15]: df_201810 = pd.read_csv('201810-fordgobike-tripdata.csv')
In [16]: df_201809 = pd.read_csv('201809-fordgobike-tripdata.csv')
In [17]: df_201808 = pd.read_csv('201808-fordgobike-tripdata.csv')
In [18]: df_201807 = pd.read_csv('201807-fordgobike-tripdata.csv')
In [19]: df_201806 = pd.read_csv('201806-fordgobike-tripdata.csv')
In [20]: df_201805 = pd.read_csv('201805-fordgobike-tripdata.csv')
In [21]: df_201804 = pd.read_csv('201804-fordgobike-tripdata.csv')
In [22]: df_201803 = pd.read_csv('201803-fordgobike-tripdata.csv')
In [23]: df_201802 = pd.read_csv('201802-fordgobike-tripdata.csv')
In [24]: df_201801 = pd.read_csv('201801-fordgobike-tripdata.csv')
In [25]: df_2017 = pd.read_csv('2017-fordgobike-tripdata.csv')

In [529]: df = df_201907.append(df_201906, sort=True)
          df = df.append(df_201905, sort=True)
          df = df.append(df_201904, sort=True)
          df = df.append(df_201903, sort=True)
          df = df.append(df_201902, sort=True)
          df = df.append(df_201901, sort=True)
          df = df.append(df_201812, sort=True)
          df = df.append(df_201811, sort=True)
          df = df.append(df_201810, sort=True)
          df = df.append(df_201809, sort=True)
          df = df.append(df_201808, sort=True)
          df = df.append(df_201807, sort=True)
          df = df.append(df_201806, sort=True)
          df = df.append(df_201805, sort=True)
          df = df.append(df_201804, sort=True)
          df = df.append(df_201803, sort=True)
          df = df.append(df_201802, sort=True)
          df = df.append(df_201801, sort=True)

In [530]: df = df.append(df_2017, sort=True)

```

4 3. Let's get some data for Muni and Bart station locations

I am assuming that the proximity to public transit makes the bike docking stations more popular.

4.1 3.1 Muni stops

```
In [28]: muni_stops = pd.read_csv('sfmta_transit/stops.txt')
muni_stops.head()
```

```
Out [28]:
```

	stop_lat	stop_code	stop_lon	stop_url	stop_id	stop_desc	\
0	37.792357	14026	-122.421010	NaN	4026	NaN	
1	37.793826	14027	-122.409591	NaN	4027	NaN	
2	37.793653	14024	-122.410823	NaN	4024	NaN	
3	37.794682	14025	-122.402770	NaN	4025	NaN	
4	37.792526	14022	-122.419589	NaN	4022	NaN	

	stop_name	location_type	zone_id
0	Clay St & Polk St	0	NaN
1	Clay St & Powell St	0	NaN
2	Clay St & Mason St	0	NaN
3	Clay St & Montgomery St	0	NaN
4	Clay St & Larkin St	0	NaN

4.2 3.2 BART stops

```
In [29]: bart= [[-122.27145,37.803768],[-122.419694,37.765062],[-122.268602,37.80835],\
[-122.418143,37.75247],[-122.270062,37.852803],[-122.447506,37.721585],\
[-122.126514,37.696924],[-122.075602,37.690746],[-122.414123,37.779732],\
[-122.196869,37.753661],[-122.466233,37.684638],[-122.029095,37.973737],\
[-122.469081,37.706121],[-122.268133,37.870104],[-122.316794,37.925086],\
[-121.899179,37.701687],[-122.39702,37.792874],[-121.976608,37.557465],\
[-122.224175,37.774836],[-122.433817,37.733064],[-122.087018,37.669723],\
[-122.12463,37.893176],[-122.26518,37.797027],[-122.26704,37.829065],\
[-122.386702,37.600271],[-122.401066,37.789405],[-122.28344,37.873967],\
[-122.024653,38.003193],[-122.212191,37.713238],[-122.183791,37.878361],\
[-122.056012,37.928468],[-121.945154,38.018914],[-122.298904,37.902632],\
[-122.407974,37.784471],[-122.353099,37.936853],[-122.251371,37.844702],\
[-122.160844,37.721947],[-122.416287,37.637761],[-122.392409,37.615966],\
[-122.057189,37.634375],[-122.44396,37.664245],[-122.017388,37.59063],\
[-121.939313,37.502171],[-122.067527,37.905522],[-121.92824,37.699756],\
[-122.29514,37.804872],[-121.7799352782858,37.9952478246996],\
[-121.8889731954402,38.01681081863409]]
```

```
In [30]: bart = pd.DataFrame(bart)
bart.columns = ['long','lat']
bart.head()
```

```
Out [30]:
```

	long	lat
0	-122.271450	37.803768

```

1 -122.419694  37.765062
2 -122.268602  37.808350
3 -122.418143  37.752470
4 -122.270062  37.852803

```

4.3 3.3 Calculating the distance of each stop from BART and Muni

```
In [481]: from geopy.distance import geodesic
```

```

def distance_calc(origin, destination):

    return geodesic(origin, destination).miles

```

```
In [32]: def bart_dist(df):
        station = (df.start_station_latitude, df.start_station_longitude)
        bart_df = list(zip(bart.lat, bart.long))

        distances = [distance_calc(station, entry) for entry in bart_df]
        return min(distances)

```

```
In [33]: def muni_dist(df):
        station = (df.start_station_latitude, df.start_station_longitude)
        muni_df = list(zip(muni_stops.stop_lat, muni_stops.stop_lon))

        distances = [distance_calc(station, entry) for entry in muni_df]
        return min(distances)

```

4.3.1 We need to filter for only San Francisco because proximity to BART and Muni is only valid here

Our original dataset includes the East Bay and San Jose, but we just want to look at San Francisco. We want to be West of -122.368535 (less than) and North of 37.702170 (greater than)

```
In [181]: df = df[df['start_station_latitude'] > 37.702170 ]
```

```
In [182]: df = df[df['start_station_longitude'] < -122.368535]
```

4.3.2 Dataframe of just the station id's, coordinates, and distances to BART and Muni

We can join this to other dataframes later, it avoids recalculating distances.

```
In [36]: coords_only = df.groupby('start_station_id').agg({'start_station_latitude': 'min', 'start_station_longitude': 'min'})
```

```
In [482]: coords_only['min_bart_dist'] = coords_only.apply(bart_dist, axis=1)
        coords_only['min_muni_dist'] = coords_only.apply(muni_dist, axis=1)

```

5 4. Let's look at some time data

```
In [531]: df['start_time'] = pd.to_datetime(df['start_time'])
```

```
In [532]: df.set_index('start_time', inplace=True)
```

5.1 4.1 Weekly sampling

```
In [533]: weekly = df.groupby('start_station_id').resample('W').count()
```

```
In [534]: weekly = weekly[['bike_id']]
```

```
In [535]: weekly.columns = ['trip_count']
```

```
In [536]: weekly.head()
```

```
Out[536]:
```

		trip_count
start_station_id	start_time	
3.0	2017-07-02	111
	2017-07-09	191
	2017-07-16	218
	2017-07-23	346
	2017-07-30	290

```
In [537]: weekly.reset_index(inplace=True)
```

```
In [538]: weekly['start_time'] = pd.to_datetime(weekly['start_time'])
```

```
In [539]: weekly['month'] = weekly.start_time.dt.month  
weekly['year'] = weekly.start_time.dt.year
```

Now joining the coordinate and distance data back to our weekly trip counts.

```
In [540]: joined_weekly = weekly.join(coords_only, on='start_station_id', how='left')
```

```
In [541]: joined_weekly.reset_index().head()
```

```
Out[541]:
```

	index	start_station_id	start_time	trip_count	month	year	\
0	0	3.0	2017-07-02	111	7	2017	
1	1	3.0	2017-07-09	191	7	2017	
2	2	3.0	2017-07-16	218	7	2017	
3	3	3.0	2017-07-23	346	7	2017	
4	4	3.0	2017-07-30	290	7	2017	

	start_station_latitude	start_station_longitude	min_bart_dist	\
0	37.786375	-122.404904	0.213256	
1	37.786375	-122.404904	0.213256	
2	37.786375	-122.404904	0.213256	
3	37.786375	-122.404904	0.213256	
4	37.786375	-122.404904	0.213256	

	min_muni_dist
0	0.014029
1	0.014029
2	0.014029
3	0.014029
4	0.014029

5.2 4.2 Plotting weekly trips

Here let's average all the trips per station over the entire timeframe so we can put them on a map with one colored marker.

```
In [542]: weekly_markers = joined_weekly.groupby('start_station_id')['trip_count'].mean()
```

```
In [543]: weekly_markers = pd.DataFrame(weekly_markers)
```

```
In [544]: weekly_markers['normalized_trip_count'] = weekly_markers.trip_count / weekly_markers.
```

```
In [545]: weekly_markers.normalized_trip_count.describe()
```

```
Out[545]: count      388.000000
          mean         1.000000
          std         1.116897
          min         0.009584
          25%         0.257491
          50%         0.657130
          75%         1.416971
          max         6.946962
          Name: normalized_trip_count, dtype: float64
```

```
In [546]: weekly_markers['marker_color'] = weekly_markers['normalized_trip_count'].apply(color)
```

```
In [547]: weekly_markers.head()
```

```
Out[547]:
```

	start_station_id	trip_count	normalized_trip_count	marker_color
3.0	574.590909	5.506844	darkred	
4.0	120.600000	1.155823	red	
5.0	474.763636	4.550106	darkred	
6.0	587.927273	5.634659	darkred	
7.0	167.027778	1.600784	darkred	

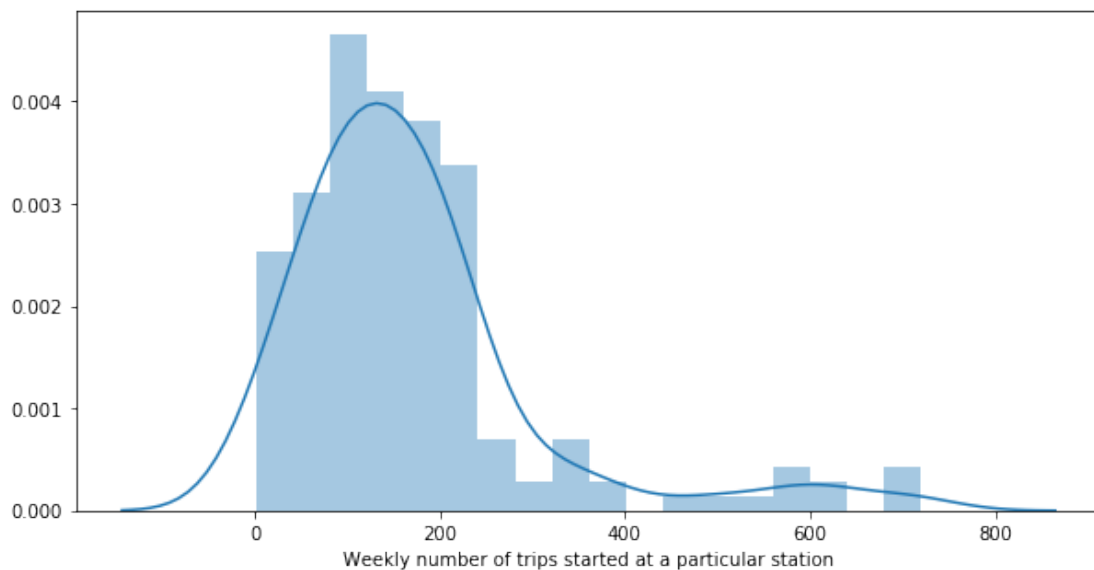
```
In [432]: def color_selector(size):
          """This function assigns a color to station markers on the map.
          Based on the normalized number of trips from that station."""

          if size < 0.5:
              color = 'yellow'
          elif size > 0.5 and size < 1:
              color = 'orange'
          elif size > 1 and size < 1.5:
              color = 'red'
          elif size > 1.5:
              color = 'darkred'
          else:
              color = 'black'
          return color
```


5.3 4.3 Weekly Trip Count Histogram

```
In [548]: import numpy as np
plt.figure(figsize=(10, 5))
sns.distplot(markers_joined['trip_count'])
plt.xlabel('Weekly number of trips started at a particular station')
```

```
Out[548]: Text(0.5, 0, 'Weekly number of trips started at a particular station')
```



5.4 4.4 Daily Sampling

Let's also look at daily sampling. Perhaps knowing the day of the week can help us make a better model.

```
In [188]: daily = df.groupby('start_station_id').resample('D').count()
```

```
In [189]: daily = daily[['bike_id']]
daily.columns = ['trip_count']
```

```
In [190]: daily.reset_index(inplace=True)
```

```
In [191]: daily.head()
```

```
Out[191]:
```

	start_station_id	start_time	trip_count
0	3.0	2017-06-29	22
1	3.0	2017-06-30	23
2	3.0	2017-07-01	27
3	3.0	2017-07-02	39
4	3.0	2017-07-03	28

```
In [192]: daily['start_time'] = pd.to_datetime(daily['start_time'])
```

```
In [193]: daily['month'] = daily.start_time.dt.month
daily['day'] = daily.start_time.dt.weekday_name
daily['year'] = daily.start_time.dt.year
```

```
In [194]: daily.head(1)
```

```
Out[194]:
```

	start_station_id	start_time	trip_count	month	day	year
0	3.0	2017-06-29	22	6	Thursday	2017

Now, joining the station coordinates and distances to transit back to the time data.

```
In [195]: daily_coords = daily.join(coords_only, on='start_station_id', how='left')
```

```
In [196]: daily_coords.head()
```

```
Out[196]:
```

	start_station_id	start_time	trip_count	month	day	year	\
0	3.0	2017-06-29	22	6	Thursday	2017	
1	3.0	2017-06-30	23	6	Friday	2017	
2	3.0	2017-07-01	27	7	Saturday	2017	
3	3.0	2017-07-02	39	7	Sunday	2017	
4	3.0	2017-07-03	28	7	Monday	2017	

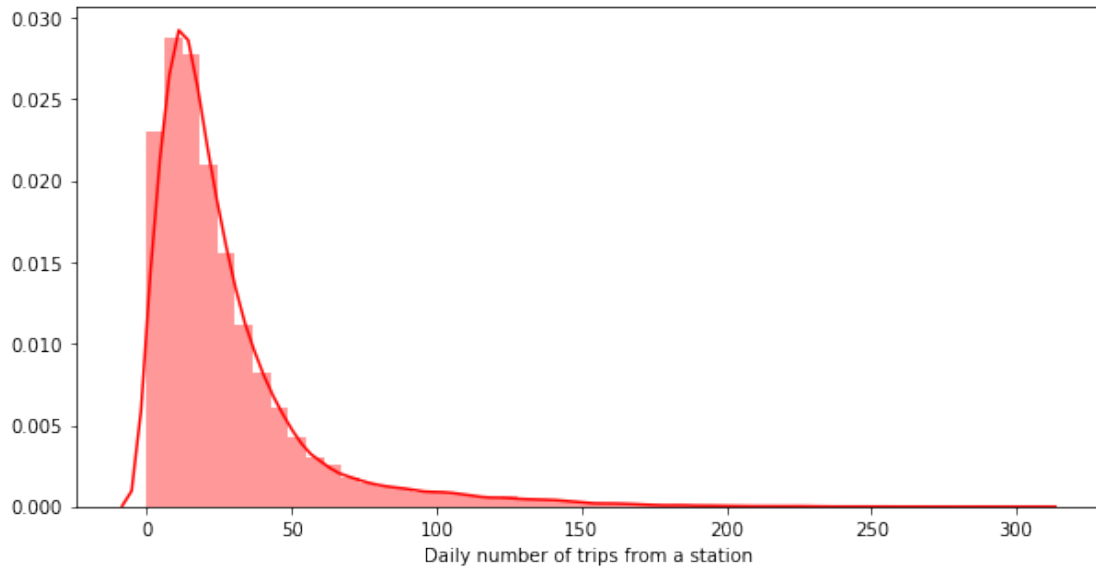
	start_station_latitude	start_station_longitude	min_bart_dist	\
0	37.786375	-122.404904	0.342934	
1	37.786375	-122.404904	0.342934	
2	37.786375	-122.404904	0.342934	
3	37.786375	-122.404904	0.342934	
4	37.786375	-122.404904	0.342934	

	min_muni_dist
0	0.022524
1	0.022524
2	0.022524
3	0.022524
4	0.022524

5.5 4.5 Daily Trip Count Histogram

```
In [197]: plt.figure(figsize=(10,5))
sns.distplot(daily_coords['trip_count'], color='red')
plt.xlabel('Daily number of trips from a station')
```

```
Out[197]: Text(0.5, 0, 'Daily number of trips from a station')
```



6 5. Let's get into some machine learning

Let's build a model that takes into account the station coordinates, distance from public transit, and some time variable to predict station popularity.

It turns out that the weekly data was superior to daily data for training the model. That is what is shown in this section.

The transformer below, will help us select which columns of the dataframe we want to feed into which part of the model.

```
In [140]: import numpy as np
          from sklearn import base

          class ColumnSelectTransformer(base.BaseEstimator, base.TransformerMixin):

              def __init__(self, columns):
                  self.columns = columns

              def fit(self, X, y=None):
                  return self

              def transform(self, X):
                  return X[self.columns].values
```

6.1 5.1 Let's separate out the label (y value) data and separate out a test set

```
In [277]: usage = joined_weekly['trip_count'].values
          data = joined_weekly
```

```
In [278]: from sklearn.utils import shuffle

         data, usage = shuffle(data, usage)

In [279]: from sklearn.model_selection import train_test_split

         X_train, X_test, y_train, y_test = train_test_split(data, usage, test_size=0.1, rand
```

6.2 5.2 Models that we'll be using:

- **K-nearest-neighbors:** The principle behind nearest neighbor methods is to find a predefined number of training samples closest in distance to the new point, and predict the label from these. Neighbors-based regression can be used in cases where the data labels are continuous rather than discrete variables. The label assigned to a query point is computed based on the mean of the labels of its nearest neighbors.
- **Ridge regression:** This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm.
- **Random forest regression:** A random forest is a meta estimator that fits a number of classifying decision trees on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting

```
In [201]: from sklearn.neighbors import KNeighborsRegressor
         from sklearn.linear_model import Ridge
         from sklearn.ensemble import RandomForestRegressor

         linreg = linear_model.Ridge()
         knn = KNeighborsRegressor()
         tree = RandomForestRegressor()
```

6.3 5.3 This is for clustering the coordinates:

```
In [280]: from sklearn.pipeline import Pipeline

         coord_pipe = Pipeline([('cst', ColumnSelectTransformer(['start_station_latitude', 's
                                ('knn', KNeighborsRegressor())
                                ])
```

6.3.1 Let's do some cross validation/ grid search to find the best hyperparameters:

```
In [281]: from sklearn import model_selection

         gs_coord = model_selection.GridSearchCV(
             coord_pipe,
             {"knn__n_neighbors": range(10, 100)},
             cv=5)
         gs_coord.fit(X_train, y_train)
         print (gs_coord.best_params_)

{'knn__n_neighbors': 37}
```

```
In [282]: gs_coord.score(X_test, y_test)
```

```
Out[282]: 0.7081739085071528
```

6.4 5.4 This is for the distance information:

```
In [514]: dist_pipe = Pipeline([('cst', ColumnSelectTransformer(['min_bart_dist', 'min_muni_di
                                ('tree', RandomForestRegressor())
                                ]))
```

```
In [515]: gs_dist = model_selection.GridSearchCV(
            dist_pipe,
            {"tree__n_estimators": [90, 140, 150, 160, 170]},
            cv=5)
            gs_dist.fit(X_train, y_train)
            print (gs_dist.best_params_)
```

```
{'tree__n_estimators': 150}
```

```
In [516]: gs_dist.score(X_test, y_test)
```

```
Out[516]: 0.7217002316330802
```

6.5 5.5 This is for the time series information:

```
In [523]: from sklearn.preprocessing import OneHotEncoder
```

```
time_pipe = Pipeline([
                                ('cst', ColumnSelectTransformer(['year', 'month'])),
                                ('ohe', OneHotEncoder(categories='auto')),
                                ('tree', RandomForestRegressor())
                                ])
```

```
In [524]: gs_time = model_selection.GridSearchCV(
            time_pipe,
            {"tree__n_estimators": [90, 140, 150, 160, 170]},
            cv=5)
            gs_time.fit(X_train, y_train)
            print (gs_time.best_params_)
```

```
{'tree__n_estimators': 160}
```

```
In [525]: gs_time.score(X_test, y_test)
```

```
Out[525]: 0.05905515969991437
```

While this prediction score of 6% may seem insignificant, when the time aspect of the prediction is removed the impact is actually closer to a 10% loss in accuracy. The month and year data is worth keeping.

6.6 5.6 Let's combine the time and spatial data into our final model:

Our previous estimators need to be converted to transformers. They will all be combined in a feature union and fed into a new pipe with a linear estimator to weight the 3 models appropriately.

```
In [152]: class EstimatorTransformer(base.BaseEstimator, base.TransformerMixin):
```

```
    def __init__(self, estimator):
        self.estimator = estimator

    def fit(self, X, y):
        self.estimator.fit(X,y)
        return self

    def transform(self, X):
        result = self.estimator.predict(X)
        mid = np.array(result)
        final = mid.reshape(-1,1)
        return final
```

```
In [302]: time_trans = EstimatorTransformer(gs_time)
          dist_trans = EstimatorTransformer(gs_dist)
          coord_trans = EstimatorTransformer(gs_coord)
```

```
In [506]: from sklearn.pipeline import FeatureUnion

          union = FeatureUnion([('time', time_trans),
                                ('space', space_trans),
                                ('coords', coord_trans)
                                ])
```

6.6.1 Here is the full model:

```
In [507]: full_model_pipe = Pipeline([('union', union), ('ridge', Ridge(alpha=10))])
          full_model_pipe.fit(X_train, y_train)
```

```
Out[507]: Pipeline(memory=None,
                    steps=[('union',
                            FeatureUnion(n_jobs=None,
                                           transformer_list=[('time',
                                                                EstimatorTransformer(estimator=Grid)
```

```

transformer_weights=None, verbose=False)),
('ridge',
 Ridge(alpha=10, copy_X=True, fit_intercept=True, max_iter=None,
        normalize=False, random_state=None, solver='auto',
        tol=0.001))),
verbose=False)

```

```
In [508]: full_model_pipe.score(X_test, y_test)
```

```
Out[508]: 0.8163848240868449
```

6.7 5.6 Let's put everything we've done above into one function:

This function will take coordinates for a proposed station and output a prediction.

```

In [526]: def coordinates_to_predictions(coordinates):
    """Takes in a list of coordinates, makes dataframe,
    calculates distances to public transit, predicts the avg number of trips per week,
    normalizes the prediction, prints out a prediction, and assigns a color to
    the marker for placing on the map.
    """

    ##### Making sure that we are in SF ###
    if (coordinates[0] <= 37.692174) or (coordinates[0] >= 37.807073):
        print('Your chosen address is not in San Francisco. Please enter a valid address')
        return None

    if (coordinates[1] >= -122.351810) or (coordinates[1] <= -122.528495):
        print('Your chosen address is not in San Francisco. Please enter a valid address')
        return None

    ##### Making a dataframe #####
    """This will fill in info for a hypothetical full year. We will average over all
    the predictions at the end."""

    d_f = pd.DataFrame()
    month = list(range(1,13))
    year = [2019]*12
    d_f['year'] = year

```

```

d_f['month'] = month
d_f['start_station_latitude'] = coordinates[0]
d_f['start_station_longitude'] = coordinates[1]

##### Calculating distance to public transit...#####
d_f['min_muni_dist'] = d_f.apply(muni_dist, axis=1)
d_f['min_bart_dist'] = d_f.apply(bart_dist, axis=1)

##### Predicting number of trips per week...#####
prediction = full_model_pipe.predict(d_f)
d_f['prediction'] = pd.Series(prediction)

weekly_trip_mean = weekly_markers.trip_count.mean()
d_f['normalized_trip_count'] = d_f.prediction/weekly_trip_mean

##### Printing out stuff for the user #####
print("The predicted number of trips per week, starting at this location, is:", )
print("The normalized weekly trip count is: ", d_f.normalized_trip_count.mean())

if d_f.normalized_trip_count.mean() < 1:
    print('The trip count is less than average, this may not be a good location.')
else:
    print('The trip count is more than average, this may be a good location.')

### A Dataframe of averages will be returned #####
d_f_avg = pd.DataFrame()
d_f_avg['start_station_latitude'] = [coordinates[0]]
d_f_avg['start_station_longitude'] = coordinates[1]
d_f_avg['trip_count'] = d_f.prediction.mean()
d_f_avg['normalized_trip_count'] = d_f.normalized_trip_count.mean()
d_f_avg['marker_color'] = 'blue'
d_f_avg['min_muni_dist'] = d_f.min_muni_dist.mean()
d_f_avg['min_bart_dist'] = d_f.min_bart_dist.mean()

return d_f_avg

```

7 6. Using Prophet to look at a time series of one station

A slight side-track looking at a time series for just one station and using Prophet to predict future demand. The analysis will look at the San Francisco Ferry Building.

```
In [433]: from fbprophet import Prophet
```

```
In [434]: stations = df.groupby('start_station_name')
```

```
In [435]: ferry_building = stations.get_group('San Francisco Ferry Building (Harry Bridges Plaza)')
```

```
In [436]: ferry_time = ferry_building.groupby('date')['date'].count()
```



```

In [437]: ferry_time = pd.DataFrame(ferry_time)
          ferry_time.head()
          ferry_time['DS'] = ferry_time.index

In [438]: ferry_time.columns = ['y', 'ds']

In [439]: ferry_time.head()

Out[439]:
           y      ds
date
2017-06-28  47  2017-06-28
2017-06-29  86  2017-06-29
2017-06-30  76  2017-06-30
2017-07-01  54  2017-07-01
2017-07-02  57  2017-07-02

In [489]: m = Prophet(daily_seasonality=False)
          m.fit(ferry_time)

Out[489]: <fbprophet.forecaster.Prophet at 0x16e559518>

In [490]: future = m.make_future_dataframe(periods=365)


In [491]: forecast = m.predict(future)
          forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()

Out[491]:
           ds      yhat  yhat_lower  yhat_upper
1124 2020-07-26  34.218068    3.320309   66.697036
1125 2020-07-27  88.174987   53.355814  122.837820
1126 2020-07-28  96.632684   64.533279  126.674147
1127 2020-07-29  90.853974   57.605654  127.383945
1128 2020-07-30  96.245663   64.764603  128.532360


In [492]: fig1 = m.plot(forecast)
          plt.xlabel('Date', fontsize=16)
          plt.ylabel('Number of Trips', fontsize=16)
          plt.title('Trips started at the Ferry building', fontsize=20)

Out[492]: Text(0.5, 1.0, 'Trips started at the Ferry building')

```



```
In [493]: fig2 = m.plot_components(forecast)
```



This nicely captures the seasonal trend in ride numbers, where there is a drop-off in November-January. It also shows the drop-off in rides on the weekend as compared to weekdays, indicating that commuters are the primary users.